

Supplementary Materials: An Efficient Approach for Signature Profiling in Genomic Data through Variable-Length k -mers

Chelsea J.-T. Ju, Jyun-Yu Jiang, Ruirui Li, Zeyu Li, and Wei Wang

SI. PROOF OF PROPOSITION 1

Proof. Given the prefix length i and the number of possible characters c , there are c^i possible prefixes in total. Assuming the characters are uniformly distributed, the probability that a particular prefix exists in n signatures is:

$$1 - \left(1 - \frac{1}{c^i}\right)^n.$$

Therefore, the expected number of collided prefixes is:

$$n - c^i \left(1 - \left(1 - \frac{1}{c^i}\right)^n\right),$$

and the expected number of prefixes without any collision is:

$$n - \left(n - c^i \left(1 - \left(1 - \frac{1}{c^i}\right)^n\right)\right) = c^i \left(1 - \left(1 - \frac{1}{c^i}\right)^n\right) = c^i \left(1 - \left(\frac{c^i - 1}{c^i}\right)^n\right).$$

However, the expected number above includes the cases that fail before reaching the i -th character. Hence, the expected number of these cases should be deducted from the above number. Finally, the expected number of signatures that fail to find their length- i prefixes along the trie during its insertion is:

$$c^i \left(1 - \left(\frac{c^i - 1}{c^i}\right)^n\right) - c^{i-1} \left(1 - \left(\frac{c^{i-1} - 1}{c^{i-1}}\right)^n\right).$$

□

SII. PROOF OF PROPOSITION 2

From Proposition 1, the expected number of node for prefix length i in n signatures is $c^i \left(1 - \left(\frac{c^i - 1}{c^i}\right)^n\right)$. Intuitively, summing all possible prefix lengths up to the length of signature m , the expected number of trie nodes is $\sum_{i=1}^m c^i \left(1 - \left(\frac{c^i - 1}{c^i}\right)^n\right)$. We provide the proper proof below.

Proof. Denote the expected number of signatures that fail to find their length- i prefixes on the trie during its insertion as $f(i)$. Given the length of signatures m , each signature that fails to find the length- i prefix along the trie during its insertion will result in the addition of $m - i + 1$ nodes. Based on Proposition 1, the expected number of nodes in the trie is:

$$\begin{aligned} \sum_{i=1}^m (m - i + 1) \cdot f(i) &= \sum_{i=1}^m (m - i + 1) \left[c^i \left(1 - \left(\frac{c^i - 1}{c^i}\right)^n\right) - c^{i-1} \left(1 - \left(\frac{c^{i-1} - 1}{c^{i-1}}\right)^n\right) \right] \\ &= \sum_{i=1}^m (m - i + 1) \left[c^i \left(1 - \left(\frac{c^i - 1}{c^i}\right)^n\right) \right] - \sum_{i=1}^{m-1} (m - i) \left[c^i \left(1 - \left(\frac{c^i - 1}{c^i}\right)^n\right) \right] \\ &= \sum_{i=1}^{m-1} [(m - i + 1) - (m - i)] \left[c^i \left(1 - \left(\frac{c^i - 1}{c^i}\right)^n\right) \right] + c^m \left(1 - \left(\frac{c^m - 1}{c^m}\right)^n\right) \\ &= \sum_{i=1}^m c^i \left(1 - \left(\frac{c^i - 1}{c^i}\right)^n\right). \end{aligned}$$

□

SIII. PROOF OF PROPOSITION 3

Proof. Suppose that the number of signatures to be added into a trie is extremely large. The expected number of nodes with c possible characters is:

$$\lim_{n \rightarrow \infty} \sum_{i=1}^m c^i \left(1 - \left(\frac{c^i - 1}{c^i} \right)^n \right) = \sum_{i=1}^m c^i = \frac{c(c^m - 1)}{c - 1} = \frac{c^{m+1} - c}{c - 1}.$$

For the plain AC, there are four possible characters, i.e., A, C, G and T. Hence, the expected number of its nodes N_A is:

$$N_A = \frac{4^{m+1} - 4}{4 - 1} = \frac{(2^2)^{m+1} - 4}{3} = \frac{2^{2m+2} - 4}{3} = \frac{4(2^{2m} - 1)}{3}.$$

For the thinned automaton, there are two possible characters, i.e., 0 and 1. Hence, the expected number of its nodes N_T is:

$$N_T = \frac{2^{m+1} - 2}{2 - 1} = 2(2^m - 1).$$

Finally, we compute the ratio of the expected number of two approaches as follows:

$$\frac{N_T}{N_A} = \frac{2(2^m - 1)}{\frac{4(2^{2m} - 1)}{3}} = \frac{3}{2} \cdot \frac{2^m - 1}{2^{2m} - 1} = \frac{3}{2} \cdot \frac{2^m - 1}{(2^m + 1)(2^m - 1)} = \frac{3}{2} \cdot \frac{1}{2^m + 1}.$$

□

SIV. RELATED WORKS AND SOFTWARE CONFIGURATIONS

Existing k -mer counters index the reads into a compact and searchable structure, such as a hash table, a burst trie, or a compact suffix array. The occurrences of a specific k -mer can be retrieved by querying these data structures. Most of these counters are designed to process reads with fixed-size k -mers; several of them restrict the choice of k to fall within a threshold. These algorithms can be adapted to count k -mers of different sizes by repeating the process with different k 's. Here, we discuss these approaches.

A. Thread-Safe Shared Memory Hashing

Jellyfish [23] exploits the CAS (compare-and-swap) assembly instruction to update a memory location in a multi-threaded environment, and uses the “quotienting technique” and bit-packed data structure to reduce wasted memory. It provides a function (`--if=kmerfile`) to count only a list of specific k -mers with the same k . The counting step is repeated for different k . We use the following two command lines to perform counting:

```
$ jellyfish count -m k -t 1 -s 100M -C --if=kmerfile readfile -o countfile
$ jellyfish dump -c countfile >> countresult
```

Squeakr [27] builds an off-the-shelf data structure based on counting quotient filter (CQF). It maintains both global and local CQFs to facilitate updates of each thread. We first count the approximate occurrences of k -mers in reads, and then query the occurrences based on a list of specific k -mers. This process is repeated for different k . Note that the exact mode cannot handle small k where all k -mers appear in the data.

```
$ squeakr count -k k -s 20 -t 1 -o k_tmp readfile
$ squeakr query -f k_tmp -q kmerfile -o countfile
```

B. Disk-Based Hashing

Disk-based hashing reduces memory with complementary disk space. In general, this approach splits k -mers into bins, and stores them in files. Each bin is then loaded into the memory for counting.

DSK [4] divides k -mers using a specific hash function based on the targeted memory and disk space. To account for arbitrary k -mer lengths, we compile the source code of DSK based on the range of the k -mer sizes (`-DKSIZE_LIST=32`). For each experiment, we first load the k -mers into memory and determine the range of the k -mer sizes. We index reads with different k 's using the main program, and dump the result into a human readable format with `dsk2ascii`.

```
$ dsk -verbose 0 -file readfile -kmer-size k -abundance-min 0
  -out tmpdir/kmers.h5 -out-tmp tmpdir -out-compress 9
$ dsk2ascii -verbose -file tmpdir/kmers.h5 -out countfile
```

MSPKmerCounter (MSPKC) [21] proposes Minimum Substring Partitioning to reduce the memory usage of storing k -mers. Observing the fact that consecutive k -mers in a read often share a shorter substring, these consecutive k -mers can be compressed and stored in one bin. It is recommended to index reads with an odd number k less than 64. The software contains three functions to be run in sequence: `partition`, `count`, and `dump`. The `partition` step divides data using minimum substring

partitioning, the count step computes the frequencies of existing k -mers, and the dump step converts the results into human readable format. This sequel is repeated for each k .

```
$ java -jar Partition.jar -in readfile -k k -L readlen -t 1
$ java -jar Count32.jar -k k -t 1
$ java -jar Dump64.jar -k k
```

KMC3 [18] scans reads one block at a time and uses a number of splitter threads to process the blocks. It leverages the concept of minimizer to further reduce disk usage. We use the main program from KMC3 to count the k -mer of all sizes seen in the list, with one k at a time. We set the `-cs` parameter to 4294967295 to ensure all of the frequently occurred k -mers are included.

```
$ kmc -t1 -srl -kK -ci0 -cs4294967295 -sfl -spl readfile countfile tmpdir
```

C. Probabilistic Hashing

Probabilistic hashing approaches avoid counting the k -mers that are likely to arise from sequencing errors.

BFCCounter [25] uses Bloom filter to identify all k -mers that are present more frequently than a threshold with a low false positive rate. The algorithm scans read data in two passes. The count function of BFCCounter requires an estimation of the number of k -mers. We use *KmerStream* [24] to pre-compute the k -mer statistics in reads. We use the dump function to convert the results into human readable format to extract the frequencies.

```
$ KmerStream -k [list of ks] -o numK_file -t 1 readfile --tsv
$ BFCCounter count -k k -n numK -t 1 -o kmers.compress readfile
$ BFCCounter dump -k k -i kmers.compress -o countfile
```

khmer [35] uses a streaming-based probabilistic data structure, CountMin Sketch [11]. The algorithm is designed to perform in-memory nucleotide sequence k -mer counting, and cannot handle k larger than 32. We use its Python wrapper script, `load-into-counting`, to perform counting, which writes a k -mer graph for each k to files. We repeat this step for different k 's. Each k -mer graph is loaded back to the memory one at a time, allowing us to query the count. We set the maximum amount of memory for data structure to be 16G as the required parameter.

```
$ khmerEnv/bin/load-into-counting.py -k k -M 16G -T 1 -q kmers.graph readfile
```

D. Suffix-Arrays

Suffix-arrays present the potential of searching arbitrary k -mers without any restriction of k on a single scan. However, constructing a suffix-array on read data can be computationally expensive.

Tallymer [19] is tailored to detect *de novo* repetitive elements ranging from 10 to 500bp in the genome. The algorithm first constructs an enhanced suffix array (`gt suffixerator`), and indexes k -mers one k at a time. We use `gt tallymer mkindex` to extract the k -mer index from the enhanced suffix-array, and use `gt tallymer search` to retrieve their counts.

```
$ gt suffixerator -dna -pl -tis -suf -lcp -v -parts 4 -db readfile -indexname tmp
$ gt tallymer mkindex -mersize k -minocc 1 -indexname kmer_index -counts -pl -esa tmp
$ gt tallymer search -output sequence counts -strand fp
    -tyr kmer_index -q kmerfile > countresult
```

MSBWT [15] compresses raw reads via a multi-string variant of Burrows-Wheeler Transform (BWT). Instead of concatenating all reads and sorting, it builds a BWT on each string and merges these multi-string BWTs through a small interleave array. The final structure allows a fast query of k -mers of arbitrary k . To index short reads, we use

```
$ msbwt cffq --uniform --compressed -p 1 tmpdir readfile
```

; to index long reads, we use

```
$ msbwt cffq -p 1 tmpdir readfile
```

. The following command line is used for querying:

```
$ msbwt massquery --rev-comp tmpdir kmerfile
```

E. Burst Tries

KCMBT [22] uses a cache efficient burst trie to store compact k -mers. The trie structure stores k -mers that share the same prefix in the same container. When a container is full, k -mers are sorted and burst. A good balance between the container size and the tree depth is essential to avoid constant sorting and bursting. As a result, KCMBT uses hundreds of trees. Unfortunately, it is limited to process k -mers with k less than 32. We first load the list of k -mers into memory. KCMBT generates binary files containing k -mers and their counts. We use `kcmbt_dump` to convert the binary data into human readable files.

```
$ kcmbt -k k -i readfile -t 1
$ kcmbt_dump 1
```

SV. SCALABILITY ANALYSIS

A. The Effect of Reads and K-mers on Single Thread

We evaluate the scalability of different approaches using the synthetic datasets. On the read level, we use 1.2 million k -mers ranging from 15-151bp (*wide*) to evaluate the effect of different read lengths and number of reads. On the k -mer level, we use 86,976,737 reads of 180bp to evaluate the effects of different batches of k -mers.

Results are shown in Table S1 and S2. The total run-time and memory of each approach are highlighted in boldface font. The run-time is further broken down into the automaton construction phase (Prep) and the read querying phase (Query) for TahcoRoll and PlainAC_Py. The two phases of MSBWT and KMC3 include indexing the reads (Prep) and querying the k -mers (Query). Read processing is performed in the querying phase of TahcoRoll and PlainAC_Py, but in the preparation phase of MSBWT and KMC3. Therefore, the run-time of querying is not on the same scale across different approaches. For Jellyfish, we use its function to count the list of k -mers directly, so the run-time cannot be split in details. Dagger (†) marks the most time efficient approach; asterisk (*) marks the most memory efficient approach.

TABLE S1: Time (Hour) and memory (GB) of profiling 1.2 million k -mers of 15-151bp on different read sets.

| Read Length | Total Reads | TahcoRoll | | | | PlainAC_Py | | | | MSBWT | | | | KMC3 | | | | Jellyfish | |
|-------------|-------------|-----------|-------|--------------|--------------|------------|-------|-------------|-------------|-------|-------|-------------|--------------|-------|-------|--------------|--------------|--------------|-------------|
| | | Prep | Query | Time | Mem | Prep | Query | Time | Mem | Prep | Query | Time | Mem | Prep | Query | Time | Mem | Time | Mem |
| 75bp | 10,128,312 | 0.006 | 0.09 | 0.10† | 3.29* | 0.02 | 0.11 | 0.13 | 6.75 | 0.40 | 0.01 | 0.41 | 5.86 | 1.49 | 0.02 | 1.51 | 3.30 | 1.83 | 4.74 |
| | 34,497,448 | 0.005 | 0.28 | 0.29† | 3.29* | 0.02 | 0.37 | 0.39 | 6.75 | 0.90 | 0.01 | 0.91 | 7.88 | 2.45 | 0.09 | 2.53 | 5.52 | 5.55 | 4.74 |
| | 97,011,938 | 0.005 | 0.78 | 0.78† | 3.29* | 0.02 | 1.04 | 1.06 | 6.75 | 3.26 | 0.02 | 1.95 | 17.57 | 5.82 | 0.68 | 6.50 | 8.00 | 15.24 | 4.74 |
| 100bp | 11,397,007 | 0.005 | 0.13 | 0.14† | 3.29* | 0.03 | 0.17 | 0.20 | 6.75 | 0.45 | 0.01 | 0.46 | 6.40 | 2.42 | 0.33 | 2.75 | 3.83 | 3.32 | 4.74 |
| | 41,054,662 | 0.005 | 0.47 | 0.48† | 3.29* | 0.02 | 0.59 | 0.61 | 6.75 | 1.29 | 0.01 | 1.30 | 11.10 | 4.81 | 0.61 | 5.42 | 7.56 | 11.25 | 4.74 |
| | 114,813,452 | 0.006 | 1.35 | 1.36† | 3.29* | 0.02 | 1.59 | 1.61 | 6.75 | 3.49 | 0.02 | 3.51 | 26.00 | 16.23 | 3.35 | 19.58 | 20.10 | 31.78 | 4.74 |
| 125bp | 10,822,319 | 0.004 | 0.15 | 0.15† | 3.29* | 0.03 | 0.19 | 0.22 | 6.75 | 0.63 | 0.01 | 0.65 | 6.83 | 2.81 | 0.81 | 3.61 | 4.15 | 5.26 | 4.74 |
| | 58,012,701 | 0.005 | 0.77 | 0.78† | 3.29* | 0.03 | 0.99 | 1.02 | 6.75 | 2.59 | 0.02 | 2.61 | 17.48 | 10.53 | 2.51 | 13.04 | 19.03 | 27.22 | 4.74 |
| | 107,375,244 | 0.005 | 1.37 | 1.38† | 3.29* | 0.02 | 1.84 | 1.87 | 6.75 | 4.56 | 0.02 | 4.58 | 29.75 | 18.46 | 3.92 | 22.37 | 34.59 | 50.41 | 4.74 |
| 150bp | 27,628,054 | 0.006 | 0.35 | 0.36† | 3.29* | 0.02 | 0.55 | 0.57 | 6.75 | 1.69 | 0.01 | 1.71 | 11.46 | 9.09 | 1.88 | 10.97 | 14.87 | 18.78 | 4.74 |
| | 57,437,772 | 0.007 | 1.20 | 1.21† | 3.29* | 0.02 | 1.20 | 1.22 | 6.75 | 3.50 | 0.02 | 3.51 | 20.31 | 17.26 | 3.98 | 21.24 | 31.10 | 36.86 | 4.74 |
| | 114,306,300 | 0.006 | 2.01 | 2.01† | 3.29* | 0.03 | 2.42 | 2.44 | 6.75 | 5.86 | 0.02 | 5.88 | 37.27 | 33.45 | 8.25 | 41.69 | 58.23 | 74.29 | 4.74 |
| 180bp | 16,197,631 | 0.006 | 0.35 | 0.35† | 3.29* | 0.03 | 0.40 | 0.43 | 6.75 | 2.43 | 0.01 | 2.45 | 9.30 | 7.45 | 1.99 | 9.44 | 14.61 | 15.51 | 4.74 |
| | 37,836,905 | 0.005 | 0.86 | 0.87† | 3.29* | 0.02 | 0.87 | 0.90 | 6.75 | 3.20 | 0.02 | 3.22 | 16.96 | 16.05 | 4.20 | 20.26 | 33.34 | 35.14 | 4.74 |

TABLE S2: Time (Hour) and memory (GB) of profiling different k -mer sets on 86,976,737 synthetic reads of 180bp.

| K-mer Batch | Total K-mers | TahcoRoll | | | | PlainAC_Py | | | | MSBWT | | | | KMC3 | | | | Jellyfish | |
|-------------------|--------------|-----------|-------|--------------|---------------|------------|-------|--------------|---------------|-------|-------|-------------|--------------|------|-------|-------------|--------------|--------------|--------------|
| | | Prep | Query | Time | Mem | Prep | Query | Time | Mem | Prep | Query | Time | Mem | Prep | Query | Time | Mem | Time | Mem |
| Small (15-31bp) | 1,200,000 | 0.0004 | 2.56 | 2.56 | 0.51* | 0.003 | 1.81 | 1.81† | 1.25 | 5.39 | 0.01 | 5.40 | 34.35 | 3.99 | 0.35 | 4.34 | 14.61 | 11.15 | 0.83 |
| | 6,000,000 | 0.002 | 4.83 | 4.83 | 2.09 | 0.02 | 2.42 | 2.44† | 5.70 | 5.39 | 0.06 | 5.46 | 34.31 | 5.39 | 0.84 | 6.23 | 14.61 | 11.11 | 0.83* |
| | 12,000,000 | 0.003 | 5.48 | 5.48 | 3.85 | 0.03 | 2.74 | 2.77† | 10.93 | 5.39 | 0.12 | 5.51 | 34.35 | 5.89 | 0.95 | 6.84 | 14.61 | 11.17 | 0.83* |
| | 24,000,000 | 0.006 | 7.22 | 7.23 | 7.13 | 0.09 | 3.11 | 3.21† | 20.93 | 5.39 | 0.23 | 5.63 | 34.35 | 5.94 | 0.96 | 6.91 | 14.61 | 11.15 | 0.83* |
| Medium (65-81bp) | 1,200,000 | 0.005 | 2.01 | 2.01† | 2.82 | 0.03 | 2.42 | 2.45 | 5.83 | 5.39 | 0.01 | 5.40 | 34.35 | 5.03 | 2.59 | 7.62 | 58.16 | 11.41 | 2.47* |
| | 6,000,000 | 0.02 | 2.47 | 2.49† | 13.49 | 0.13 | 4.77 | 4.90 | 28.59 | 5.39 | 0.06 | 5.45 | 34.35 | 4.90 | 1.91 | 6.81 | 58.16 | 11.37 | 2.47* |
| | 12,000,000 | 0.09 | 3.53 | 3.62† | 26.52 | 0.27 | 5.27 | 5.54 | 56.71 | 5.39 | 0.11 | 5.50 | 34.33 | 4.90 | 1.93 | 6.83 | 58.16 | 11.07 | 2.47* |
| | 24,000,000 | 0.16 | 4.00 | 4.16† | 52.11 | 0.75 | 5.25 | 6.00 | 112.5 | 5.39 | 0.22 | 5.61 | 34.35 | 4.87 | 1.49 | 6.37 | 58.16 | 11.36 | 2.47* |
| Large (131-151bp) | 1,200,000 | 0.02 | 2.65 | 2.67† | 6.10 | 0.06 | 2.98 | 3.04 | 12.24 | 5.39 | 0.02 | 5.41 | 34.35 | 3.51 | 2.35 | 5.87 | 67.27 | 6.66 | 4.74* |
| | 6,000,000 | 0.08 | 3.51 | 3.59† | 29.91 | 0.29 | 4.34 | 4.63 | 60.63 | 5.39 | 0.08 | 5.47 | 34.35 | 4.28 | 4.04 | 8.32 | 67.27 | 6.72 | 4.74* |
| | 12,000,000 | 0.18 | 4.37 | 4.55† | 58.43 | 0.55 | 4.98 | 5.53 | 118.97 | 5.39 | 0.16 | 5.55 | 34.33 | 5.14 | 4.50 | 9.65 | 69.19 | 8.59 | 4.38* |
| | 24,000,000 | 0.42 | 4.42 | 4.84† | 117.79 | 1.33 | 4.90 | 6.23 | 240.67 | 5.39 | 0.29 | 5.69 | 34.35 | 4.10 | 3.05 | 7.17 | 67.27 | 6.73 | 4.74* |

B. Improvement with Multiple-Threads




Fig. S1: Multi-threads evaluation of MSBWT, KMC3, Jellyfish and TahcoRoll while profiling *wide* batches of *k*-mers on 86,976,737 synthetic reads of 180bp.

TABLE S3: Time (Hour) and memory (GB) of profiling different *wide* batch of *k*-mers on 86,976,737 synthetic reads of 180bp. Dagger([†]) marks the most time efficient approach; asterisk (*) marks the most memory efficient approach.

| Total K-mers | Methods | 4-Thread (Hour) | 8-Thread (Hour) | 16-Thread (Hour) | Memory (GB) |
|--------------|-----------|-------------------|-------------------|-------------------|-------------|
| 1,200,000 | TahcoRoll | 0.70 [†] | 0.38 [†] | 0.22 [†] | 3.50* |
| | MSBWT | 2.29 | 1.83 | 1.62 | 30.83 |
| | Jellyfish | 18.29 | 9.66 | 6.10 | 4.74 |
| | KMC3 | 26.68 | 20.79 | 21.61 | 72.83 |
| 6,000,000 | TahcoRoll | 1.47 [†] | 0.77 [†] | 0.41 [†] | 16.07 |
| | MSBWT | 2.35 | 1.89 | 1.71 | 30.83 |
| | Jellyfish | 19.39 | 8.86 | 6.15 | 4.74* |
| | KMC3 | 26.55 | 22.47 | 16.65 | 72.77 |
| 12,000,000 | TahcoRoll | 1.24 [†] | 1.12 [†] | 0.64 [†] | 31.49 |
| | MSBWT | 2.42 | 1.90 | 1.76 | 30.83 |
| | Jellyfish | 18.66 | 9.94 | 6.29 | 4.74* |
| | KMC3 | 25.22 | 17.33 | 15.21 | 73.08 |
| 24,000,000 | TahcoRoll | 1.98 [†] | 1.32 [†] | 0.91 [†] | 61.86 |
| | MSBWT | 2.51 | 2.33 | 2.11 | 30.83 |
| | Jellyfish | 18.96 | 9.88 | 6.24 | 4.74* |
| | KMC3 | 22.22 | 18.74 | 15.97 | 73.08 |

SVI. EVALUATION OF BINARIZED REPRESENTATIONS

TABLE S4: Evaluation of different binarized representations. Time is reported in hour and memory is reported in gigabyte. Nucleotides can be divided into balanced or unbalanced partitions. The *p*-values are computed through paired *t*-tests on time against the default setting: $\{A, G\}, \{C, T\}$, and adjusted by Bonferroni correction.

| Mapping | 0={A, C}; 1={G, T} | | 0={A, T}; 1={C, G} | | 0={A}; 1={C, G, T} | | 0={C}; 1={A, G, T} | | 0={G}; 1={A, C, T} | | 0={T}; 1={A, C, G} | |
|-----------------|--------------------|-------|--------------------|-------|--------------------|------|--------------------|------|--------------------|------|--------------------|------|
| | Time | Mem | Time | Mem | Time | Mem | Time | Mem | Time | Mem | Time | Mem |
| SRR1293902 | 1.28 | 4.49 | 1.27 | 4.42 | 1.54 | 3.36 | 1.42 | 3.20 | 1.54 | 3.12 | 1.93 | 2.98 |
| SRR1293901 | 1.35 | 4.49 | 1.39 | 4.42 | 1.72 | 3.36 | 1.71 | 3.20 | 1.79 | 3.12 | 1.96 | 2.98 |
| GSM1254204 | 1.26 | 4.49 | 1.26 | 4.42 | 1.46 | 3.36 | 1.64 | 3.20 | 1.70 | 3.12 | 2.01 | 2.98 |
| SRR5951587 | 0.32 | 11.11 | 0.30 | 10.84 | 0.34 | 9.85 | 0.47 | 7.93 | 0.57 | 7.85 | 0.54 | 9.27 |
| SRR5951588 | 0.52 | 11.11 | 0.48 | 10.84 | 0.58 | 9.85 | 0.83 | 7.93 | 1.21 | 7.85 | 0.81 | 9.27 |
| SRR5951600 | 0.76 | 11.11 | 0.74 | 10.84 | 0.78 | 9.85 | 1.46 | 7.93 | 1.36 | 7.85 | 1.11 | 9.27 |
| <i>p</i> -value | 0.3408 | | 0.12714 | | 0.033552 | | 0.043938 | | 0.008988 | | 0.010212 | |